Operating Systems (OS) – Theory

Assignment:03

Multi-Level Feedback Queue(Simulation)

Submitted by:

Tayyaba Asif

2019-CE-5

Submitted to:

Sir Waqas Ali

Date: 25th February, 2022

Question No.01:

Run a few randomly-generated problems with just two jobs and two queues; compute the MLFQ execution trace for each. Make your life easier by limiting the length of each job and turning off I/O's

Answer:

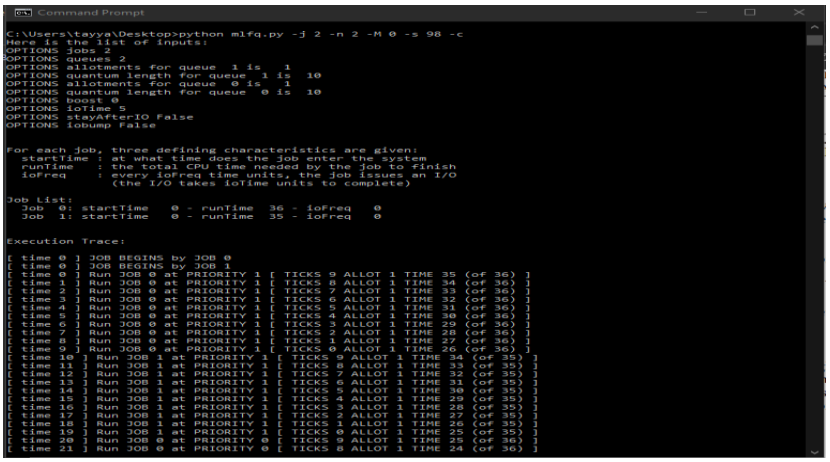

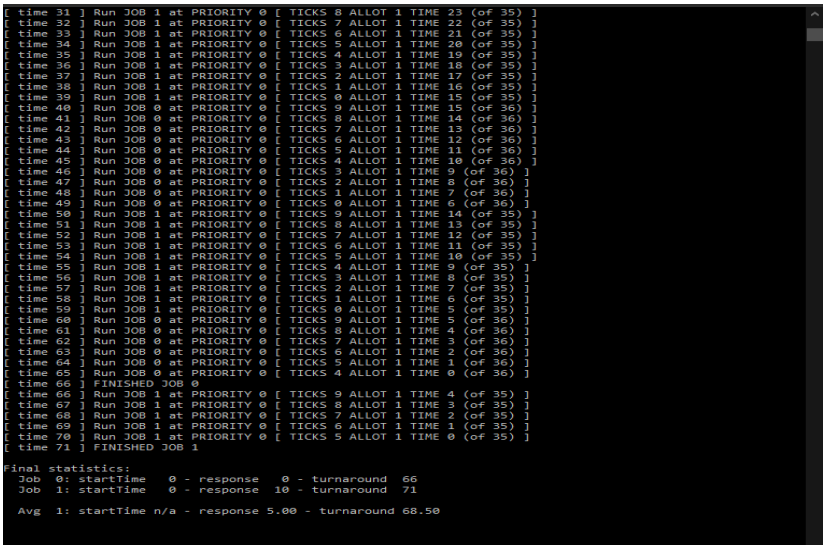

Question No.02:

How would you run the scheduler to reproduce each of the examples in the chapter?

Answer:

For figure 8.2: **python mlfq.py -l 0,200,0** this will allot 1 job to 3 queues with pre-defined quantum 10.

For figure 8.3: **python mlfq.py -l 0,200,0:100,20,0** this will allot 2 jobs to 3 queues.

For figure 8.4: **python mlfq.py -l 0,200,0:0,20,1 -i 10 -S**

For figure 8.5: **python mlfq.py -l 0,200,0:20,10,0:30,10,0:40,10,0:50,10,0:60,10,0:70,10,0:80,10,0:90,10,0:100,10,0:110,10,0:120,10,0:130,10,0:140,10,0:150,10,0:160,10,0:170,10,0:180,10,0:190,10,0:200,10,0:210,10,0:220,10,0, -B 50**

For figure 8.6: **python mlfq.py -l 0,120,0:20,100,9 -S -i 1**

For figure 8.7: **python mlfq.py -l 0,100,0:0,100,0 -Q 10,20,40**

Question No.03:

How would you configure the scheduler parameters to behave just like a round-robin scheduler?

Answer:

Jobs on the same queue are time-sliced, so the answer is to only have one queue.

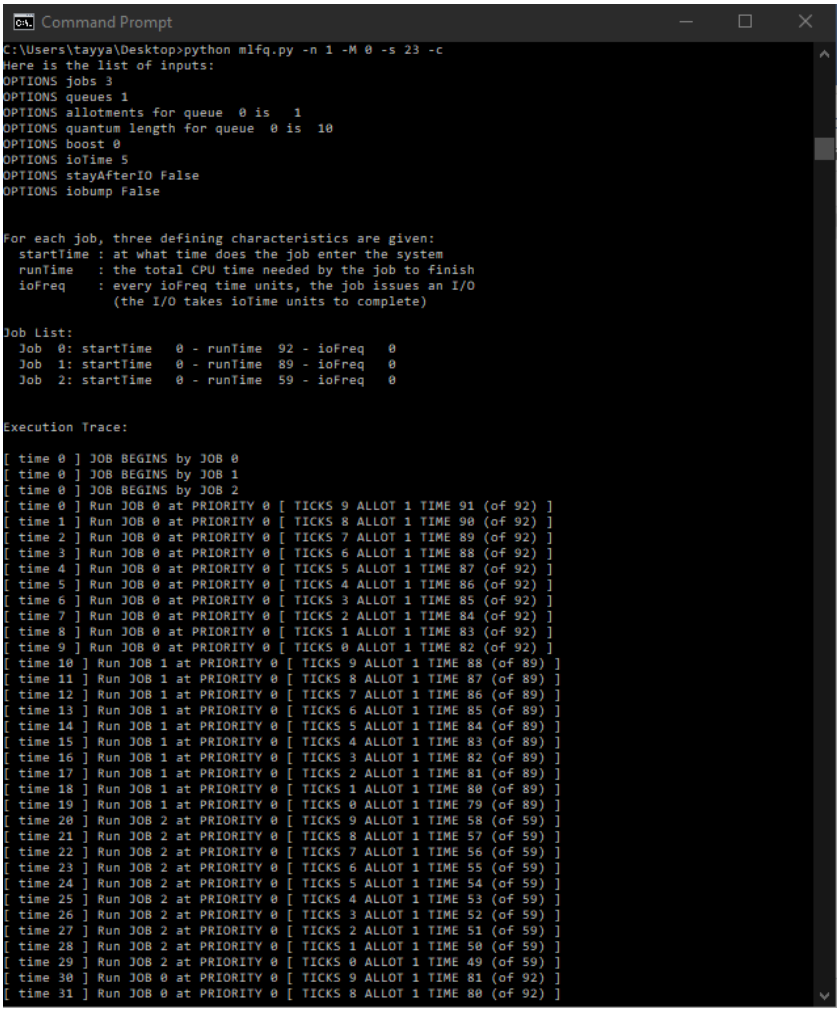

Question No.04:

Craft a workload with two jobs and scheduler parameters so that one job takes advantage of the older Rules 4a and 4b (turned on with the -S flag) to game the scheduler and obtain 99% of the CPU over a particular time interval.

Answer:

First, set the I/O time to 1 with -i, then ensure that CPU hog calls an I/O operation on the 99th click of a 100 click quantum length. Also, include the -I flag to ensure the CPU gets back control when the I/O request finishes.

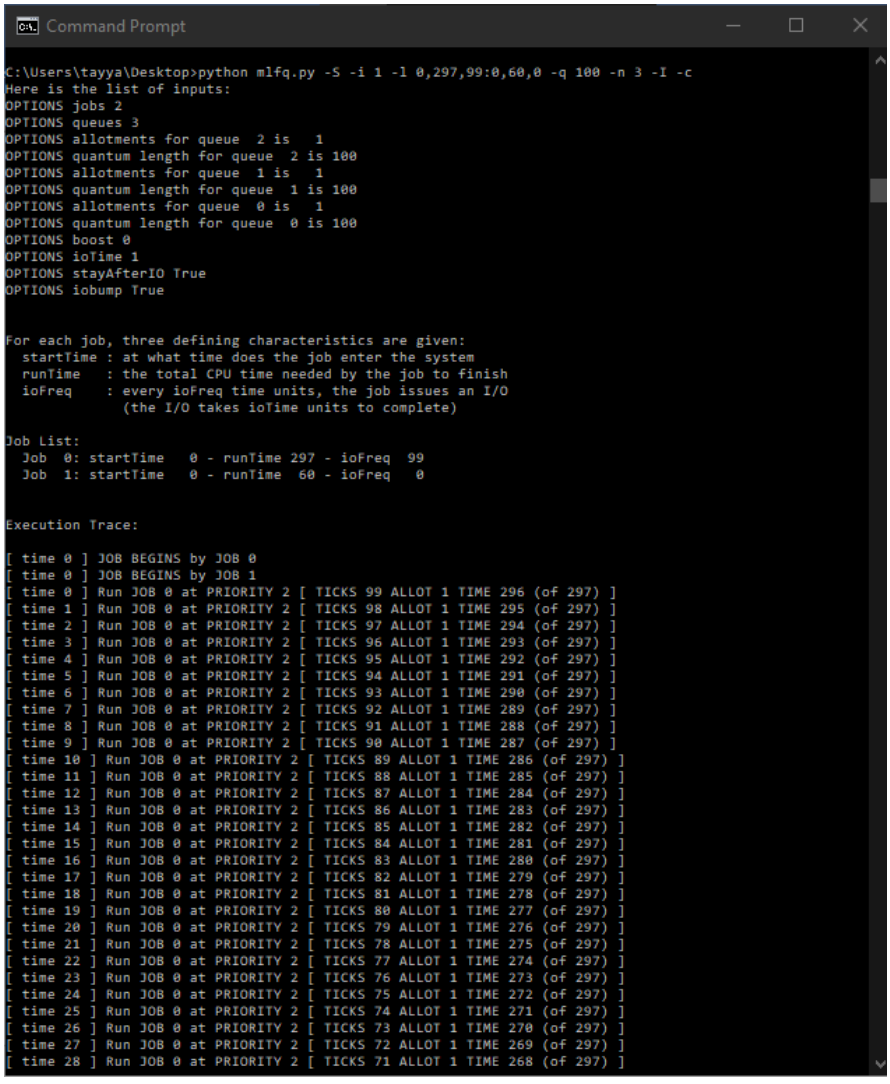

Question No.05:

Given a system with a quantum length of 10 ms in its highest queue, how often would you have to boost jobs back to the highest priority level (with the -B flag) in order to guarantee that a single long running (and potentially-starving) job gets at least 5% of the CPU?

Answer:

If a new 10ms-long job arrives every 10ms - like pictured in Figure 8.5 (left), boosting every 200 ticks ensures the long process will run 10/200 CPU ticks = 5%.